

Reinforcement Learning: Introduction

Lecture 12

Dr. Syed Maaz Shahid

20th May, 2024

Contents

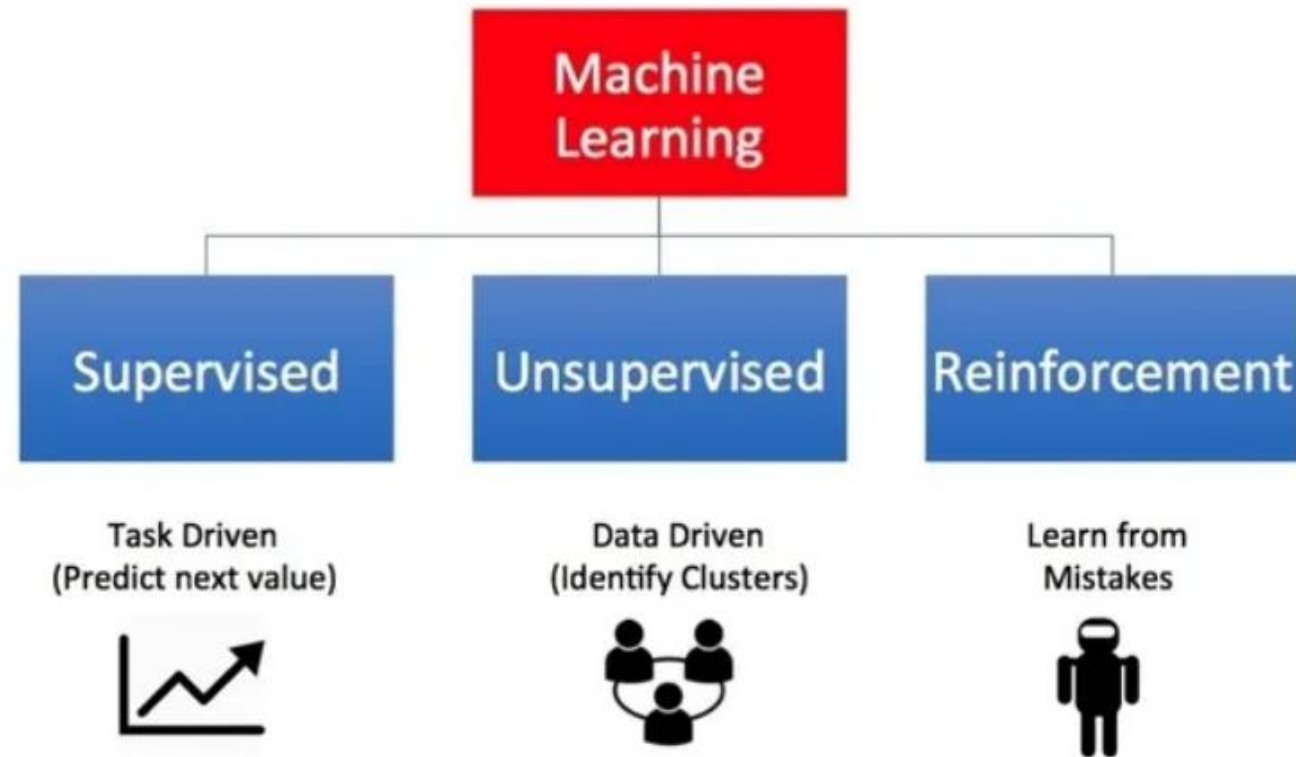
- Overview
- Elements of Reinforcement Learning
- Use Cases
- Agents and States
- Components of an RL Agent
- RL Agent Taxonomy
- Exploration and Exploitation
- Major Challenges of RL
- Multi-Agent Reinforcement Learning

Assignment

- Find a paper that uses HMM to solve a problem in your relevant field.
- Make a report and submit by **31** May, 20204.

Machine Learning

Types of Machine Learning




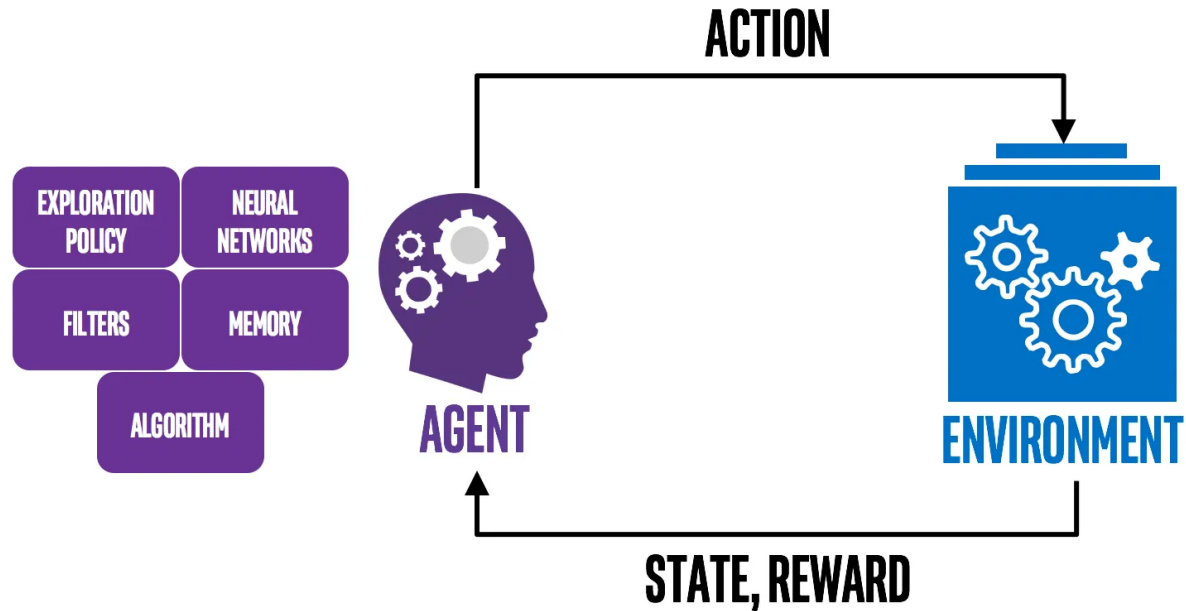
Reinforcement Learning

- In supervised machine learning, you focus on predicting what you don't know → *learning to predict*.
- Reinforcement learning (RL) is a framework for learning-based decision-making → *learning to act*.
- Reinforcement learning uses **rewards and punishments** as signals for positive and negative behavior.

Reinforcement Learning: Elements

- **Agent/ Policy:** The program you train for specific task/Method to map agent's state to actions.
- **Environment/State:** The world, real or virtual, in which the agent performs actions.
- **Action.** A move made by the agent, which causes a status change in the environment.
- **Rewards.** The evaluation of an action (positive or negative)/Feedback from the environment..


Reinforcement Learning: Elements



Reinforcement Learning: Use Case

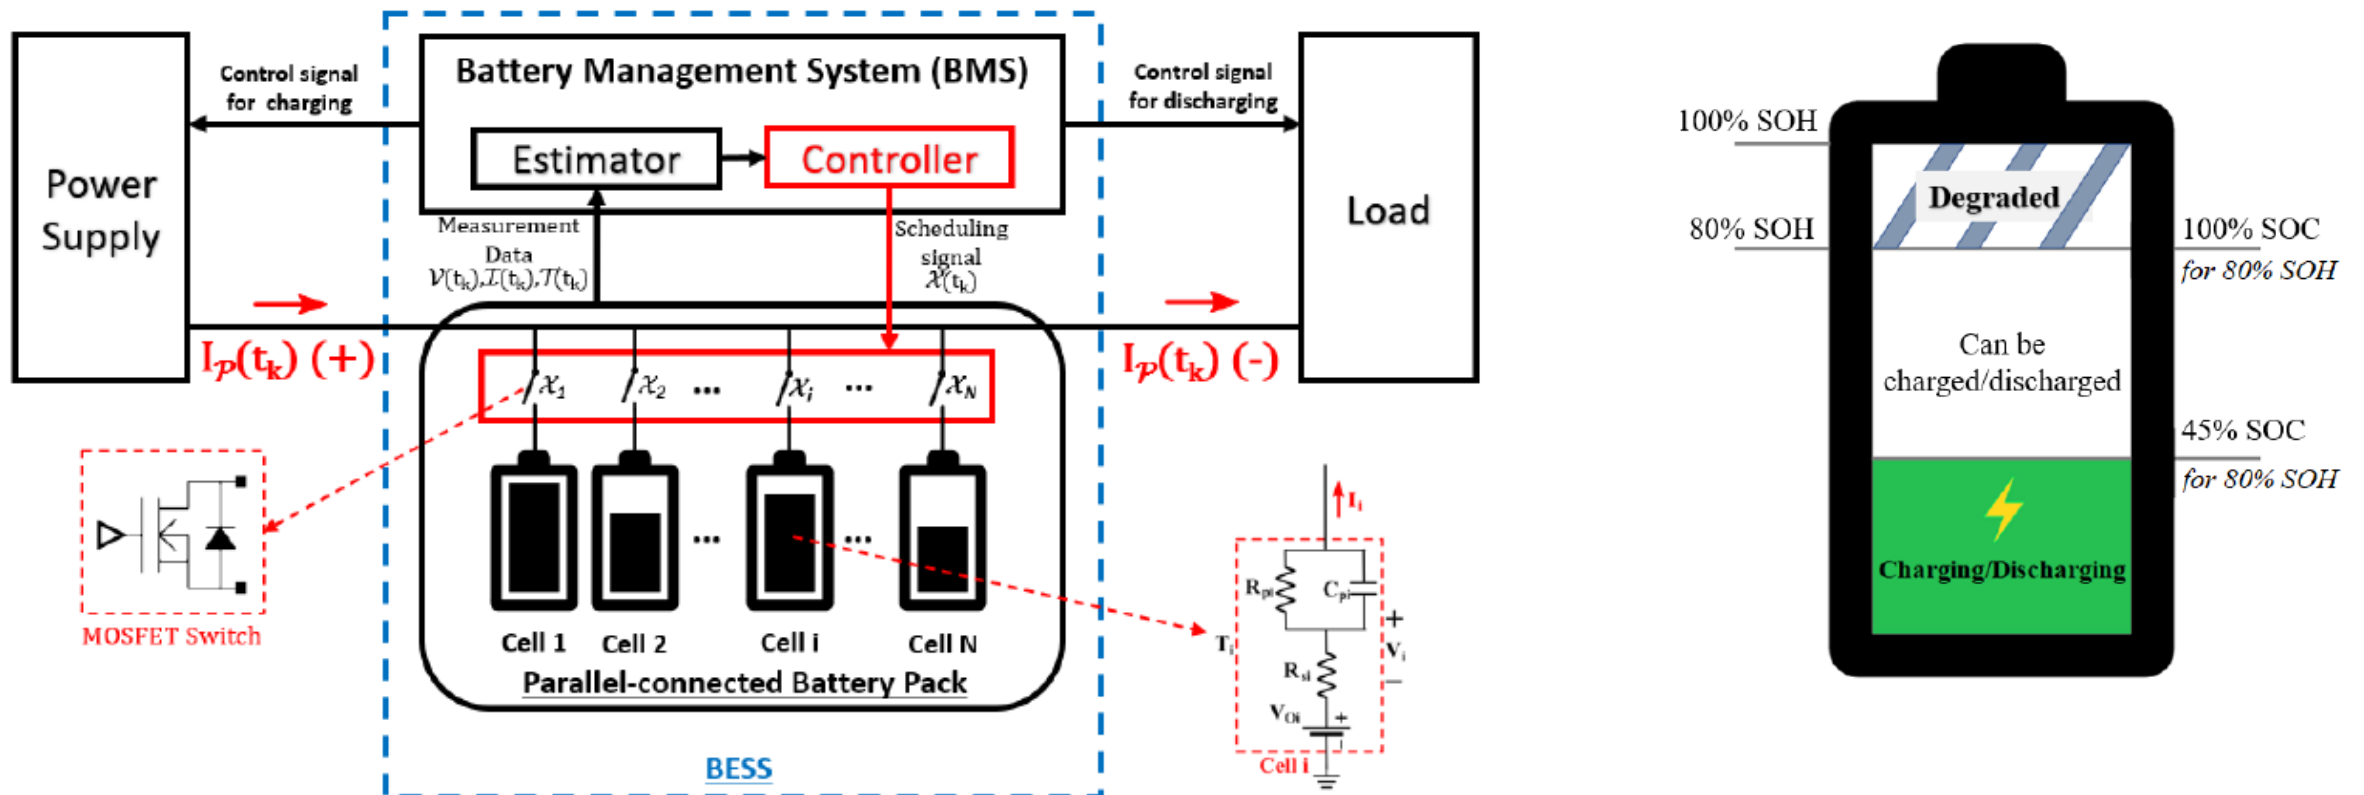
- **PacMan** game

- *Agent* → PacMan
- *Environment*: Grid world
- *Reward* for eating food & *Punishment* if gets killed by the ghost
- *State* → Location
- *Cumulative reward* → winning the game



Reinforcement Learning: Use Case

- Battery Management for Retired Electric Vehicle Batteries



Reinforcement Learning: Use Case

- Deep Q-learning playing Atari
- AlphaGo - The Movie

Differences: Supervised, Unsupervised, and Reinforcement Learning

- Static vs. Dynamic
- No explicit right answer → agent learns by trial and error
- RL requires exploration
- RL is a multiple-decision process → decision-making chain through the time in RL

Reinforcement Learning: Agent & Environment

- At each step t the agent
 - Executes action A_t
 - Receives observation O_t
 - Receives scalar reward R_t
- The environment:
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits scalar reward R_{t+1}


Reinforcement Learning: History & State

- The **history** H_t is the sequence of observations, actions, rewards
- **State** is the information used to determine what happens next.
- Formally, state is a function of the history:

$$S_t = f(H_t)$$

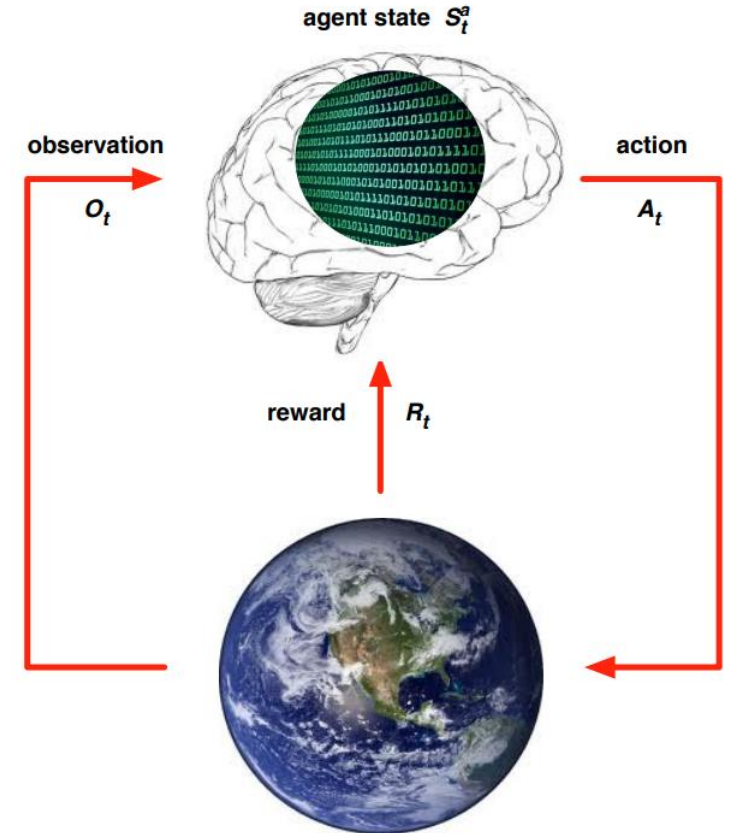
Reinforcement Learning: Environment State

- The **environment state** S_t^e is the environment's private representation.
- Usually, S_t^e is not visible to the agent.
- Even if S_t^e is visible to agent, it may contain irrelevant information



Reinforcement Learning: Agent State

- The **agent state** S_t^a is the agent's internal representation \rightarrow uses to pick the next action.
-
- It can be any function of history:
$$S_t^a = f(H_t)$$



Reinforcement Learning: Information State

- An **information state** (a.k.a. *Markov state*) contains all useful information from the history.
 - Markov chain is to assume that X_k captures all *the relevant information* for predicting the future

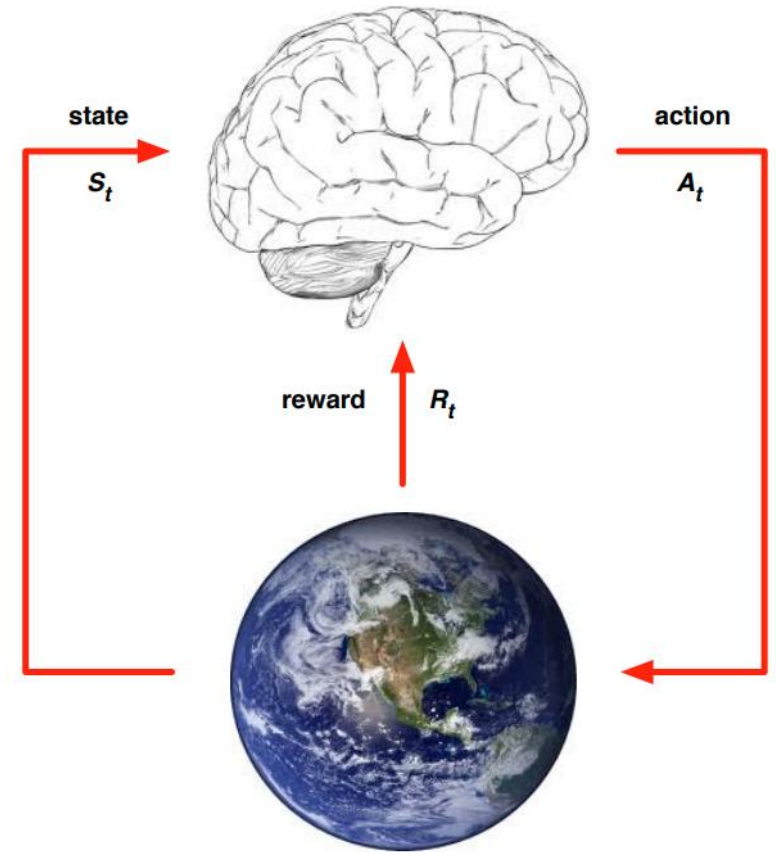
- A state S_t is Markov if and only if

$$P[S_{t+1}|S_t] = P[S_{t+1}|S_1, \dots, S_t]$$

- The future is independent of the past given the present
- Once the state is known, the history may be thrown away

Reinforcement Learning: Fully Observable Environments

- Agent directly observes environment state
$$O_t = S_t^a = S_t^e$$
- This is a Markov decision process (MDP).



Reinforcement Learning: Partially Observable Environments

- Agent indirectly observes environment

$$S_t^a \neq S_t^e$$

→ A poker playing agent only observes public cards.


- This is a partially observable Markov decision process (POMDP).
- Agent must construct its own state representation $S_t^a = H_t$

Components of an RL Agent

- **Policy:** agent's behavior function
- **Value function:** how good is each state and/or action
- **Model:** agent's representation of the environment

Policy

- A policy is the agent's behavior.
- It is a map from state to action,
 - **Deterministic** policy: $a = \pi(s) \rightarrow$ same action give a state
 - **Stochastic** policy: $\pi(a|s) = P[A_t = a|S_t = s] \rightarrow$ choose an action randomly based on the probability distribution



Value Function


- Value function $v_{\pi}(s)$ is a prediction of future reward
- Used to evaluate the goodness/badness of states
- By following a policy π , the value function is defined as
$$v_{\pi}(s) = E[R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots | S_t = s]$$
 - $0 \leq \gamma \leq 1$: *discount rate*
 - γ close to 1: rewards further in the future count more \rightarrow agent is farsighted

Model

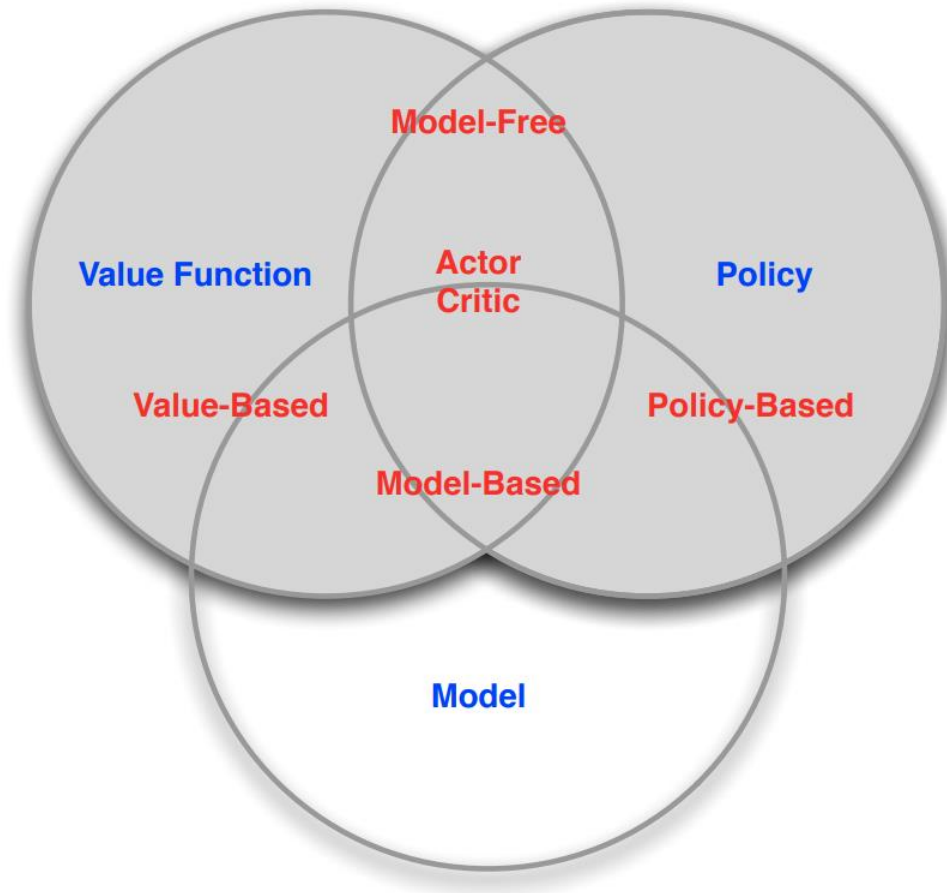
- The model describes the **environment** by a distribution over rewards and state transitions.

$$P_{s'}^a = P[S_{t+1} = s' | S_t = s, A_t = a]$$

$$R_{s'}^a = E[R_{t+1} | S_t = s, A_t = a]$$



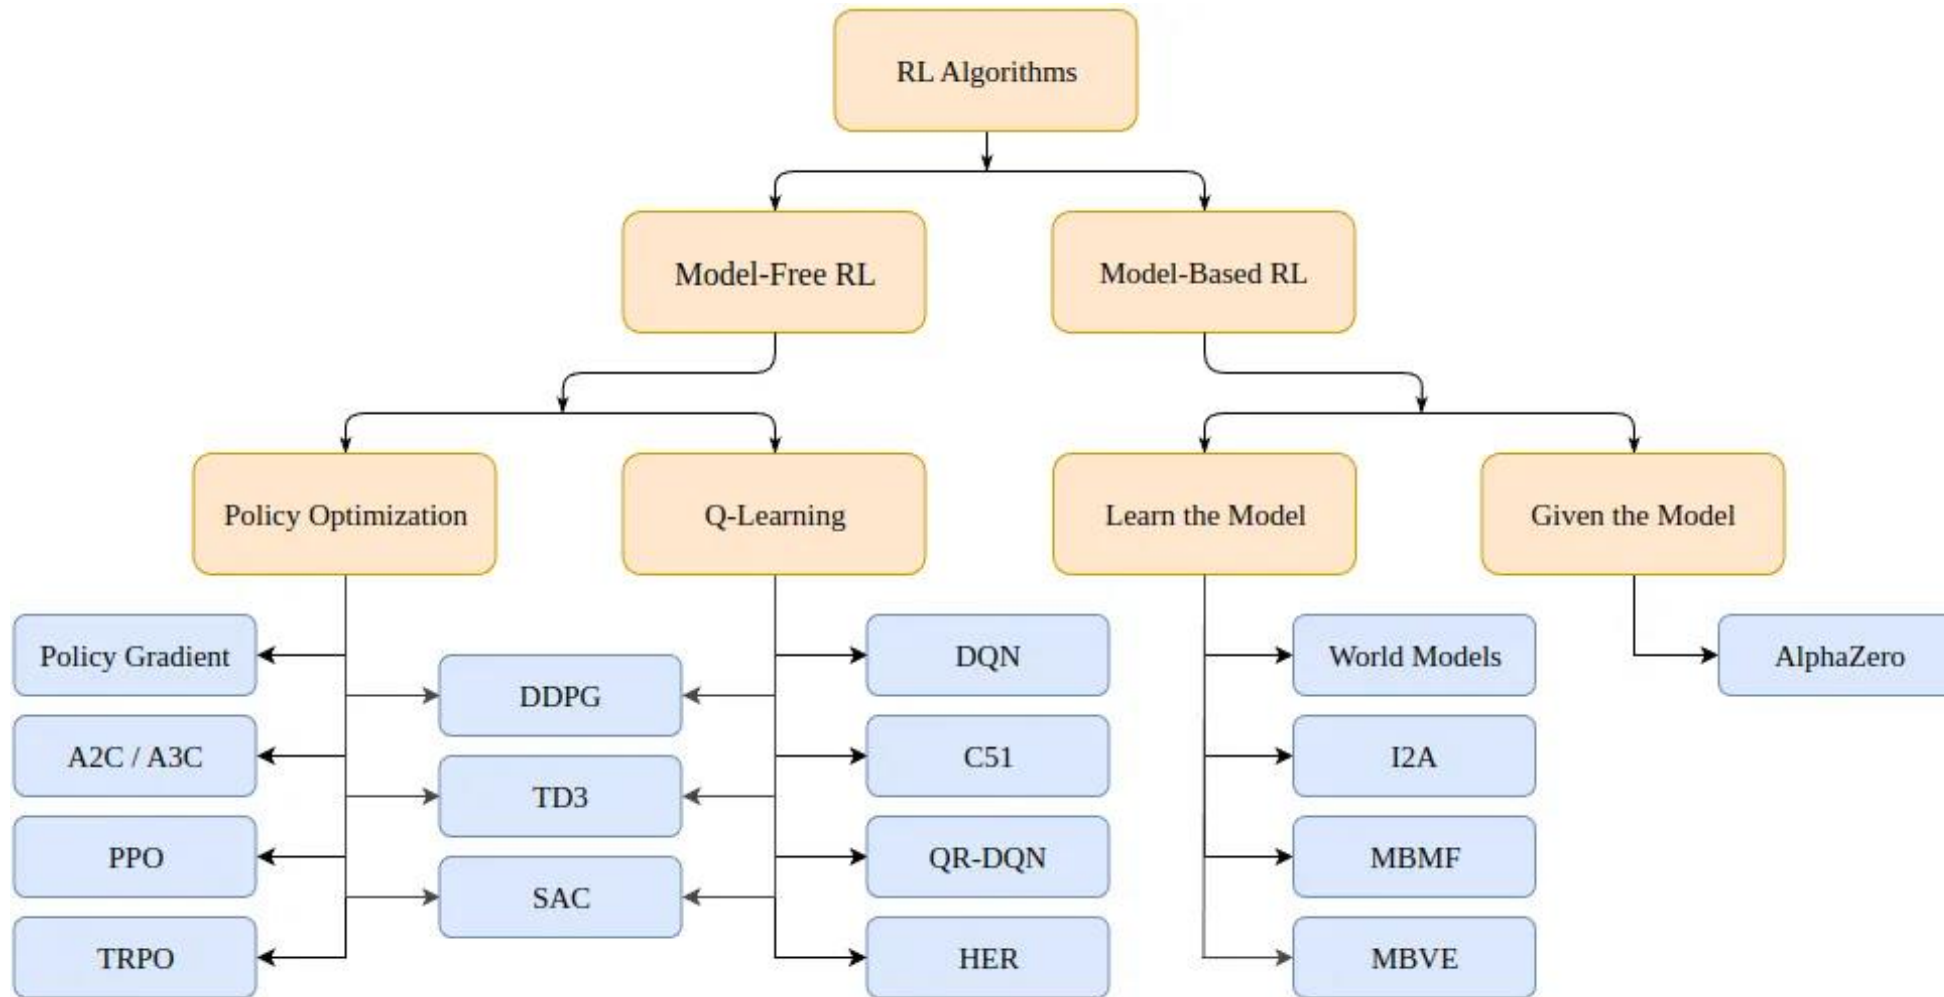
RL Agent Taxonomy



RL Agent Taxonomy

- DQN → value-based method
- A2C → policy-based method (actor) & value-based method (critic)
- Policy Optimization (PPO)

Taxonomy of Reinforcement Learning algorithms



OpenAI spinning up

Reinforcement Learning algorithms

- **Model-based** algorithms use a model of the environment.
 - Model is used to predict future states and rewards.
 - The model is either given (e.g. a chessboard) or learned.
- **Model-free** algorithms directly learn how to act for the states encountered during training
 - Which state-action pairs yield good rewards (Q-Learning).

Reinforcement Learning: Fundamental Problems

- **Learning:**

- The environment is initially unknown → which states are good or what the actions do.
- The agent interacts with the environment
- The agent improves its policy

- **Planning:**

- A model of the environment is known → Markov decision problem
- The agent performs computations with its model
- The agent improves its policy

Exploration and Exploitation

- Reinforcement learning is like **trial-and-error** learning.
 - The agent should discover a good policy from its experiences of the environment → without losing too much reward along the way.
- *Exploration*: finds more information about the environment.
- *Exploitation*: exploits known information to maximize reward.
- Interesting trade-off:
 - immediate reward (**exploitation**) vs. gaining knowledge that might enable higher future reward (**exploration**)

Exploration and Exploitation: Examples

- Restaurant Selection
 - **Exploitation**: Go to your favorite restaurant
 - **Exploration**: Try a new restaurant
- Oil Drilling
 - **Exploitation**: Drill at the best-known location
 - **Exploration**: Drill at a new location

Major Challenges of RL

- Sample efficiency
 - RL algorithms require a large amount of data and experience to learn effectively → costly and time-consuming.
- State and action spaces
 - Exponential growth of state and action spaces as problem complexity increases.
 - For example, in the game of Go, the number of possible board configurations is estimated to be 10^{170} .

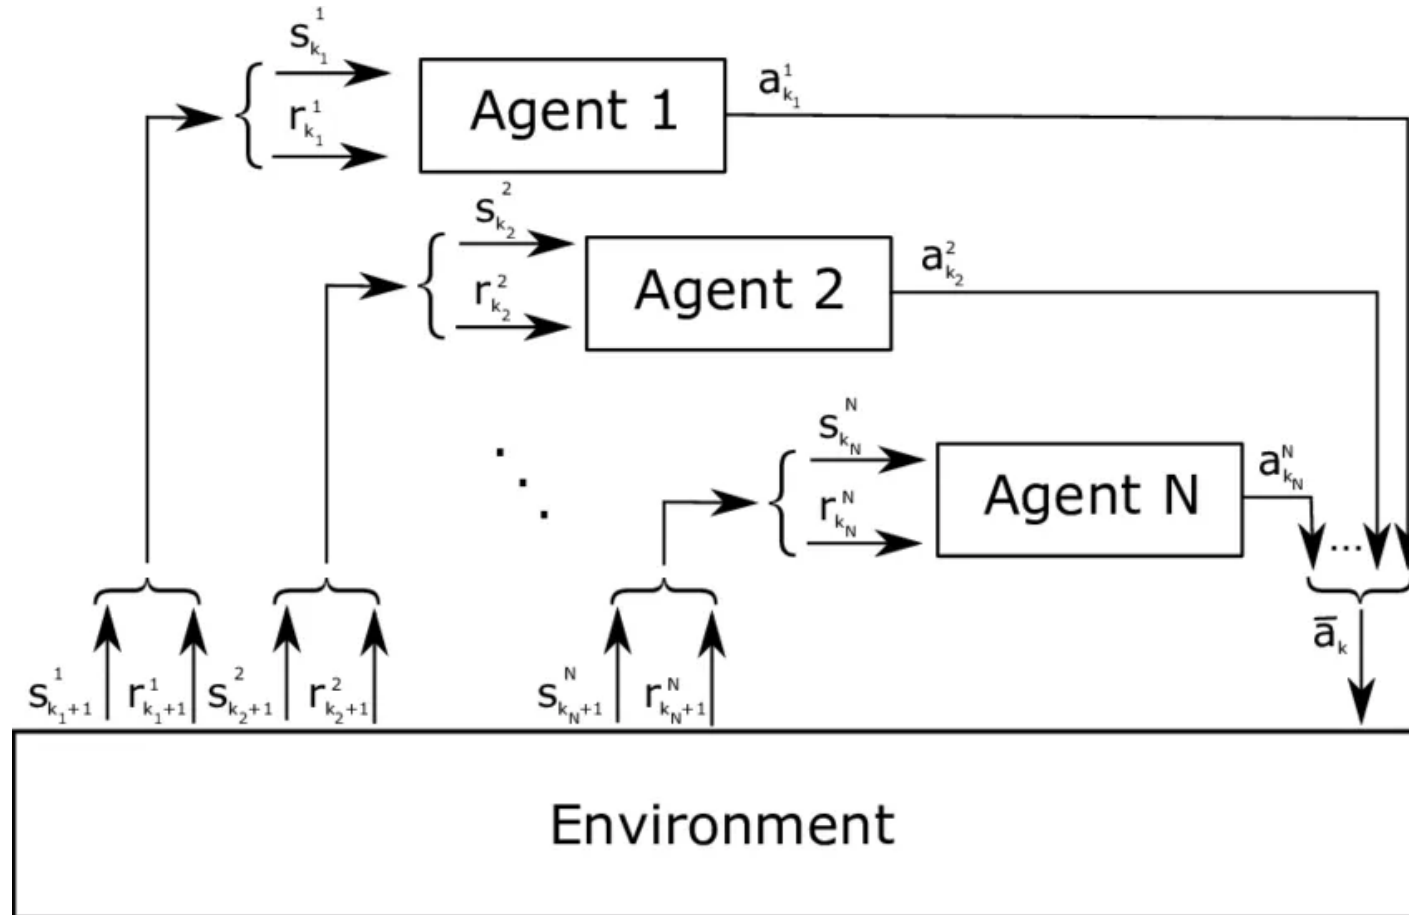
Major Challenges of RL

- Exploration and exploitation
 - Both exploration and exploitation are essential for learning, but they can also conflict with each other.

Multi-Agent Reinforcement Learning

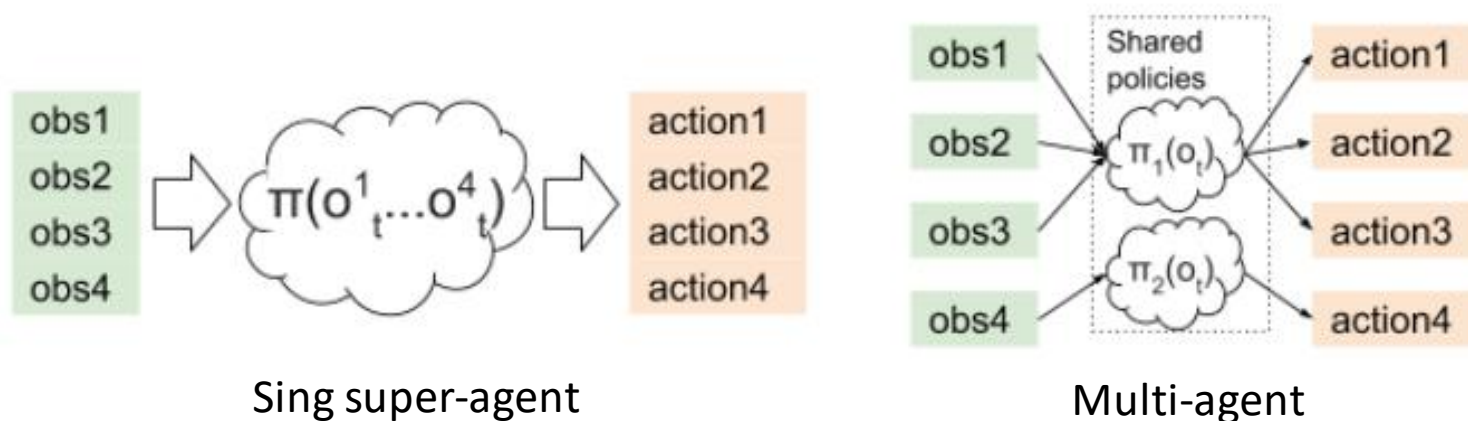
- **Vanilla reinforcement learning** is concerned with a single agent.
- Multi-agent reinforcement learning (MARL) studies how multiple agents interact in a common environment.
 - **Cooperative**: All agents working towards a common goal
 - **Competitive**: Agents competing with one another to accomplish a goal

Multi-Agent Reinforcement Learning




Why MARL?

- *A more natural decomposition of the problem:* Need to train policies for **cellular antenna tilt control**.
 - Instead of training a single *super-agent* that controls all the cellular antennas in a city → it is more natural to model each antenna as a separate agent in the environment.
- *Potential for more scalable learning*



MARL: Use case

- Dota 2: AI agents are trained to coordinate with each other to compete against humans.



MARL: Challenges

- Environment non-stationarity




Fig: Non-stationarity of environment: Initially **(a)**, the red agent learns to regulate the speed of the traffic by slowing down. However, over time the blue agents learn to bypass the red agent **(b)**, rendering the previous experiences of the red agent invalid.

References

- Reinforcement Learning: An Introduction by Richard S. Sutton and Andrew G. Barto.
- [David Silver Course on RL](#)